

Weather Forecasting Using Pyspark

Big Data Processing

Nguyen Phuoc Nguyen Phuc

Code github link: <https://github.com/WinerDeCoder/Big-Data-Weather-ForeCasting.git>



This project is inspired by: <https://github.com/andrea-gasparini/big-data-weather-forecasting>

I. Overview

This project's goal is:

- Apply the power of distributed system (Spark) into big data processing
- Preprocessing the data, transformation, aggregate using Dataframe
- A ML model to predict weather condition based on some statistics

II. Data

A. Data overview

Data is taken from [Kaggle](#) contains hourly weather measurements data of 36 cities, collected from 2012 to 2017. This 5 years of data result in approximately 45.000 measurements (for each city) of temperature, humidity, air pressure and the like.

weather_description.csv (21.86 MB)

Detail

Compact

Column

10 of 37 columns

datetime	Vancouver	Portland	San Francisco	Seattle	Los Angeles
<div><div>02/20/2017 - 05/25/2017</div><div>Count: 2,262</div></div> <div>2012-10-012017-11-30</div>	<div>sky is clear28%</div> <div>light rain12%</div> <div>Other (26879)59%</div>	<div>sky is clear26%</div> <div>light rain17%</div> <div>Other (25852)57%</div>	<div>sky is clear28%</div> <div>mist18%</div> <div>Other (24525)54%</div>	<div>sky is clear28%</div> <div>light rain16%</div> <div>Other (25290)56%</div>	<div>sky is clear28%</div> <div>light rain16%</div> <div>Other (25290)56%</div>
2012-10-01 12:00:00					
2012-10-01 13:00:00	mist	scattered clouds	light rain	sky is clear	mist
2012-10-01 14:00:00	broken clouds	scattered clouds	sky is clear	sky is clear	sky
2012-10-01 15:00:00	broken clouds	scattered clouds	sky is clear	sky is clear	sky
2012-10-01 16:00:00	broken clouds	scattered clouds	sky is clear	sky is clear	sky
2012-10-01 17:00:00	broken clouds	scattered clouds	sky is clear	sky is clear	sky
2012-10-01 18:00:00	broken clouds	scattered clouds	sky is clear	few clouds	sky

Data Explorer

Version 2 (74.69 MB)

city_attributes.csv

humidity.csv

pressure.csv

temperature.csv

weather_description.csv

wind_direction.csv

wind_speed.csv

II. Data

A. Data overview

- The data set has 7 csv files
- The city_attributes.csv are only contain city and countries name, we won't use this file as these are already in other files (however it can be used if you want to perform operation like join)
- Other 6 files, The weather_condition.csv are the weather condition we try to predict later. The rest are features we will use to predict

=> We will use only these features 'Country', 'Latitude', 'Longitude', 'humidity', 'pressure', 'temperature', 'wind_direction', 'wind_speed'

=> Our label is 'weather_condition'

B. Data preprocessing

We will process the data as following:

1. Load data on

```
weather_conditions_df = ks.read_csv('data/weather_description.csv')
humidity_df = ks.read_csv('data/humidity.csv')
pressure_df = ks.read_csv('data/pressure.csv')
temperature_df = ks.read_csv('data/temperature.csv')
city_attributes_df = ks.read_csv('data\city_attributes.csv')
wind_direction_df = ks.read_csv('data/wind_direction.csv')
wind_speed_df = ks.read_csv('data/wind_speed.csv')
```

1a. We drop some column of cities due to resource limitation of local computer (optional)

```
columns_to_drop = [ 'Charlotte', 'Miami',
                    'Pittsburgh', 'Toronto', 'Philadelphia', 'New York', 'Montreal',
                    'Boston', 'Beersheba', 'Tel Aviv District', 'Eilat', 'Haifa',
                    'Nahariyya', 'Jerusalem']

weather_conditions_df = weather_conditions_df.drop(columns_to_drop, axis=1)
humidity_df = humidity_df.drop(columns_to_drop, axis=1)
pressure_df = pressure_df.drop(columns_to_drop, axis=1)
temperature_df = temperature_df.drop(columns_to_drop, axis=1)
wind_direction_df = wind_direction_df.drop(columns_to_drop, axis=1)
wind_speed_df = wind_speed_df.drop(columns_to_drop, axis=1)
```



```

weather_conditions_melted = weather_conditions_df.melt(id_vars=['datetime'], var_name='city', value_name='weather_condition')
humidity_mel = humidity_df.melt(id_vars=['datetime'], var_name='city', value_name='humidity')
pressure_mel = pressure_df.melt(id_vars=['datetime'], var_name='city', value_name='pressure')
temperature_mel = temperature_df.melt(id_vars=['datetime'], var_name='city', value_name='temperature')
wind_direction_mel = wind_direction_df.melt(id_vars=['datetime'], var_name='city', value_name='wind_direction')
wind_speed_mel = wind_speed_df.melt(id_vars=['datetime'], var_name='city', value_name='wind_speed')

```

2. Join all tables together to form a consistence dataframe:

- In each table, we need to melt down so there no more column in each city - which is good for joining step because we will join based on “city”. After this, we should have all the table in form ['city', 'datetime', 'value'] (value is based on the feature)
- Then we join all table together on keys 'city' and 'datetime'
- Then drop records that have null value

```

join_kdf = weather_conditions_melted.join(city_attributes_df.set_index(['City']), on = 'city', how = 'left') \
.join(humidity_mel.set_index(['city', 'datetime']), on = ['city', 'datetime'], how = 'outer') \
.join(pressure_mel.set_index(['city', 'datetime']), on = ['city', 'datetime'], how = 'outer') \
.join(temperature_mel.set_index(['city', 'datetime']), on = ['city', 'datetime'], how = 'outer') \
.join(wind_direction_mel.set_index(['city', 'datetime']), on = ['city', 'datetime'], how = 'outer') \
.join(wind_speed_mel.set_index(['city', 'datetime']), on = ['city', 'datetime'], how = 'outer')

join_kdf.columns

```

```

not_null_weather_measurements_df = join_kdf.dropna()

```

3. There're a lot of weather condition but similar to each other , we will replace all of it into 6 main weather: sunny, rainy, snowy, foggy, thunderstorm, cloudy

```
weather_mapping = {  
    'sky is clear': 'sunny',  
    'overcast clouds': 'cloudy',  
    'light rain': 'rainy',  
    'broken clouds': 'cloudy',  
    'few clouds': 'cloudy',  
    'haze': 'foggy',  
    'very heavy rain': 'rainy',  
    'thunderstorm with rain': 'thunderstorm',  
    'smoke': 'foggy',  
    'scattered clouds': 'cloudy',  
    'proximity shower rain': 'rainy',  
    'fog': 'foggy',  
    'moderate rain': 'rainy',  
    'proximity thunderstorm': 'thunderstorm',  
    'light snow': 'snowy',  
    'light shower snow': 'snowy',  
    'snow': 'snowy',  
    'sleet': 'snowy',  
    'light shower sleet': 'snowy',  
    'mist': 'foggy',  
    'proximity thunderstorm': 'thunderstorm',  
    'thunderstorm with heavy rain': 'thunderstorm'  
}
```

```
not_null_weather_measurements_df['weather_condition'] = \  
not_null_weather_measurements_df['weather_condition'].replace(weather_mapping)
```

```
valid_values = ['thunderstorm', 'rainy', 'snowy', 'cloudy', 'foggy', 'sunny']  
not_null_weather_measurements_df = not_null_weather_measurements_df[not_null_weather_measurements_df['weather_condition'].isin(valid_values)]
```

4. When counting weather condition, we should get like this. However, You can see the imbalance of label
=> We will downsample all other label down to the number of min label (6833)

```
weather_condition
rainy          121854
snowy          9931
sunny          511357
cloudy         419912
thunderstorm   6833
foggy          103404
Name: count, dtype: int64
```

```
spark_df = not_null_weather_measurements_df.to_spark()

# Specific number to downsample to
specific_number = value_counts.min()

# Downsample each group
def downsample_to_specific_number(df, col_name, count):
    return df.groupby(col_name).applyInPandas(
        lambda pdf: pdf.sample(n=count, random_state=42) if len(pdf) > count else pdf,
        schema=df.schema
    )

downsampled_spark_df = downsample_to_specific_number(spark_df, 'weather_condition', specific_number)

# Convert back to Koalas DataFrame if needed
balanced_df = downsampled_spark_df.to_koalas()
```


III. Spark ML

Now we will use Pyspark for further data analysis and ML model

1. Create spark session
2. Load data from Koala_into Spark
3. Shuffle the data
4. One-hot encode 'weather_condition' to get numeric label

```
import databricks.koalas as ks
from pyspark.sql import SparkSession
from pyspark.ml.feature import OneHotEncoder, StringIndexer, VectorAssembler
from pyspark.ml.classification import DecisionTreeClassifier
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.sql import SparkSession, functions as F

spark = SparkSession.builder.appName("KoalasAndSparkML").getOrCreate()

spark_df = balanced_df.to_spark()

spark_df = spark_df.orderBy(F.rand())

# Index the weather_condition column
indexer = StringIndexer(inputCol="weather_condition", outputCol="label")
indexed_df = indexer.fit(spark_df).transform(spark_df)
```

B. Data Analysis

Before fitting this data into ML model, we should perform some analysis on it first for better result and minimize cost

1. It's necessary to view the correlation of this dataframe

a. We see the correlation of each features together

b. Then the correlation of each feature to the label

=> We do this with Correlation Matrix on all dataframe

```
from pyspark.ml.stat import Correlation
from pyspark.ml.feature import VectorAssembler

# convert to vector column first
vector_col = "corr_features"
assembler = VectorAssembler(inputCols=['humidity', 'pressure', 'temperature', 'wind_direction', 'wind_speed', 'Longitude', 'Latitude', 'label'],\
                             outputCol=vector_col)
df_vectorer = assembler.transform(indexed_df).select(vector_col)

# get correlation matrix
matrix = Correlation.corr(df_vectorer, vector_col)

matrixer = matrix.collect()[0]["pearson({})".format(vector_col)].values

import seaborn as sns
import matplotlib.pyplot as plt

x_labels = ['humidity', 'pressure', 'temperature', 'wind_direction', 'wind_speed', 'Longitude', 'Latitude', 'label']
y_labels = ['humidity', 'pressure', 'temperature', 'wind_direction', 'wind_speed', 'Longitude', 'Latitude', 'label']

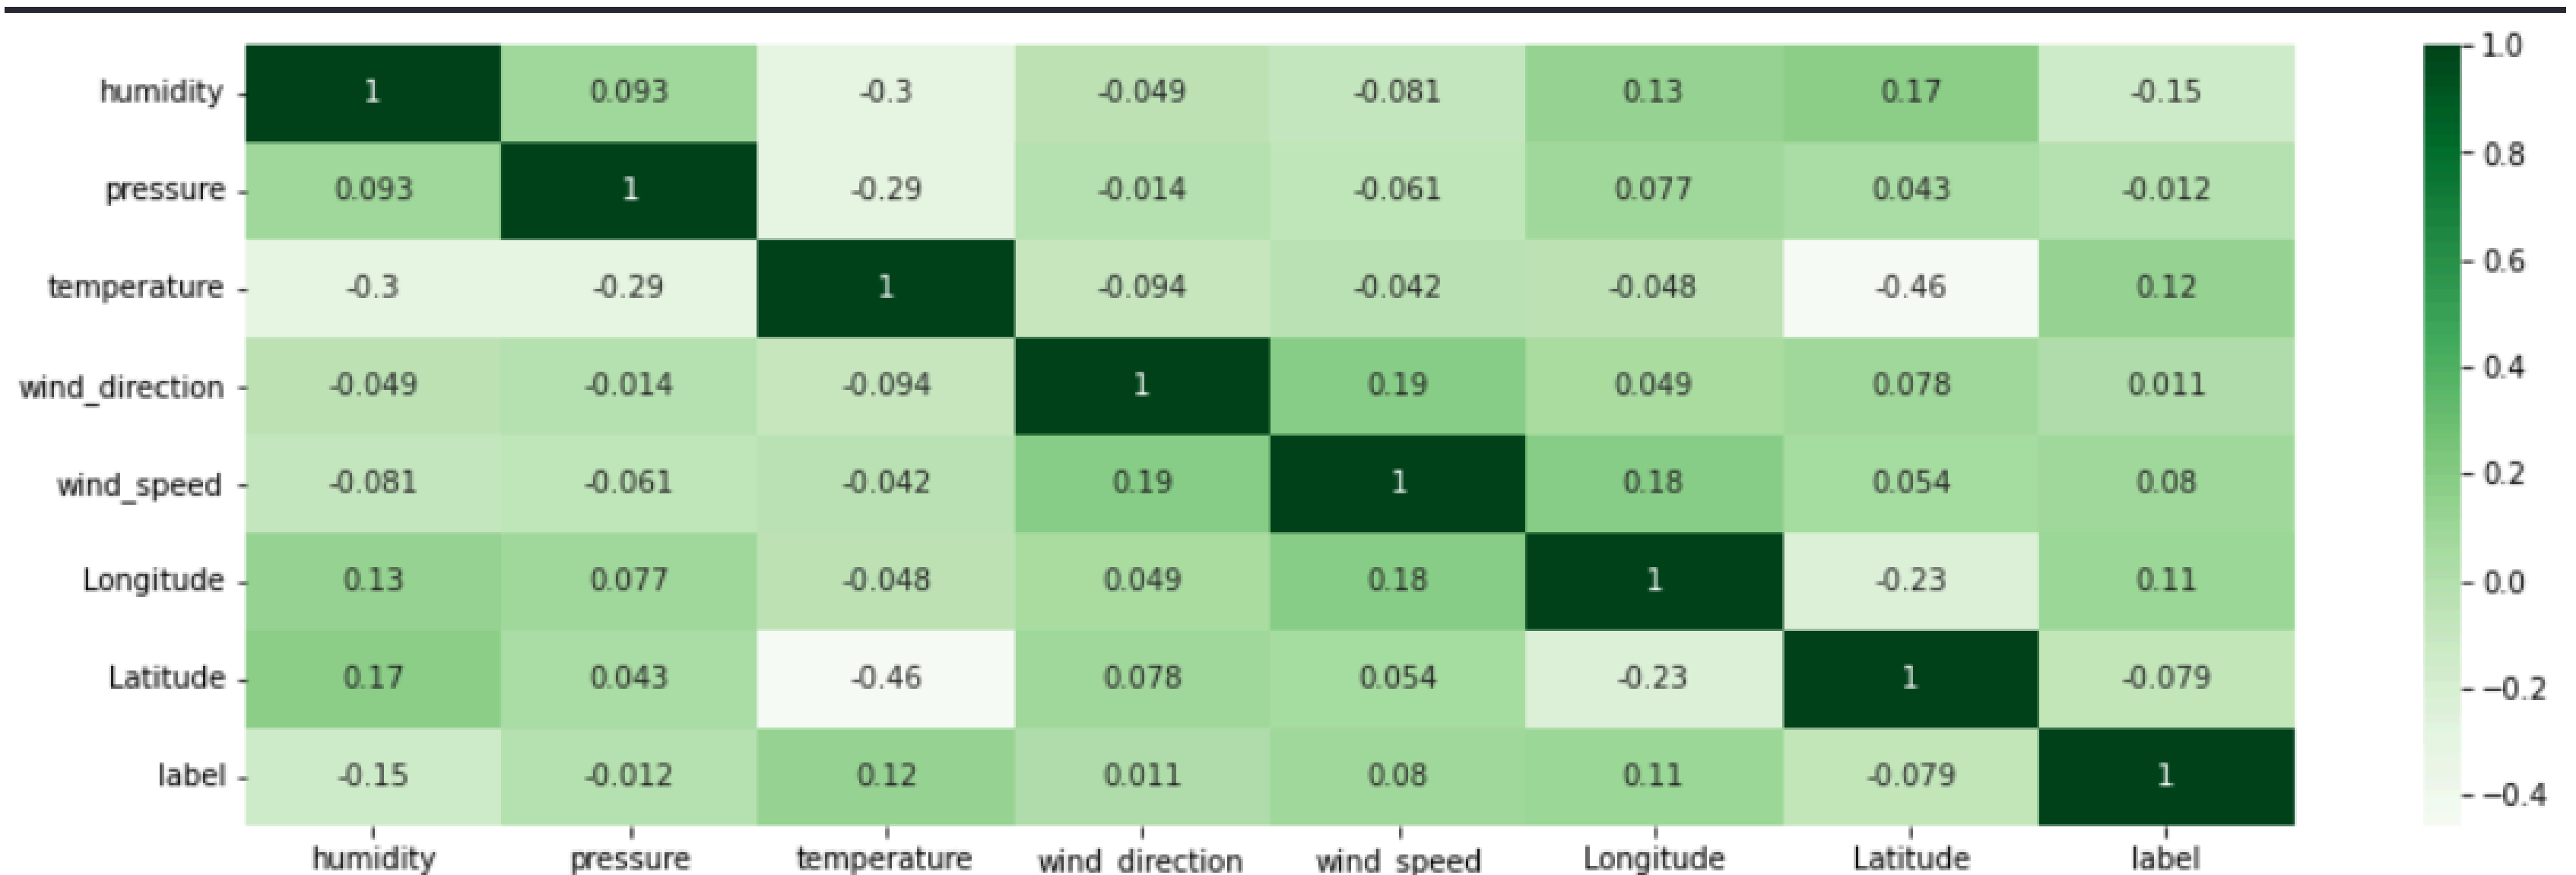
plt.figure(figsize=(16,5))
sns.heatmap(matrixer.reshape(8,8),
            xticklabels=x_labels,
            yticklabels=y_labels, cmap="Greens", annot=True)
```

2. Let see this correlation matrix:

- Basically, if the value is 0, that mean these 2 feature are completely independent. And if value in positive $0 < \text{value} \leq 1$, that mean these 2 value are increasing dependent , and $-1 \leq \text{value} < 0$ are decreasing dependent
- We want the correlation of features - label have relationship => Differ than 0. But we want correlation of feature - feature close to 0, because if 2 value is dependent, there is no sense to have both of them as feature, 1 is enough

=> As we can see

- For feature - label: only 'humidity', 'temperature', 'wind_speed', 'Longitude' and 'Latitude' seem have stronger relation ship to the label than other
- For feature - feature: Some really strong dependent are: Latitude - temperature, temperature - humidity, pressure - temperature, Latitude - Longitude



3. Feature Selection + Dimensionally Reduction:

- We observe some feature contribute more in the result ==> We only select these features.
- However, these features are have some high correlation pair => we need to dimensionally reduction the data.

For better feature into the model, lower cost but with meaningful features.

I choose to use PCA as the reduction method, From 5 feature I will down to only 3.

** Actually i think LDA is a better method here, because they are caring about feature label => It make more sense to separate these class. However, pyspark doesn't implement it yet so PCA is just fine.

```
from pyspark.ml.feature import VectorAssembler, PCA
from pyspark.ml import Pipeline

assembler = VectorAssembler(
    inputCols=['humidity', 'temperature', 'Longitude', 'wind_speed', 'Latitude'],
    outputCol='features'
)
feature_df = assembler.transform(indexed_df)

pca = PCA(k=3, inputCol='features', outputCol='pca_features')
pca_model = pca.fit(feature_df)
pca_df = pca_model.transform(feature_df)
```

III. ML model with Pyspark

For this task, to simplify, we will use Decision Tree as our Classification model (max_height = 5)

The idea of Decision Tree is :

- In each split, it will chose the feature split that maximize the information gain, try to make child node impurity. Then keep going into child node

The flow of using pyspark ML

- 1.Assembler defining “features” for ML model
- 2.PCA
- 3.Split the data into train/test with propotion 0.8/0.2
- 4.Put into the model
- 5.Testing result


```
from pyspark.ml.feature import VectorAssembler, PCA
from pyspark.ml import Pipeline

assembler = VectorAssembler(
    inputCols=['humidity', 'temperature', 'Longitude', 'wind_speed', 'Latitude'],
    outputCol='features'
)
feature_df = assembler.transform(indexed_df)

pca = PCA(k=4, inputCol='features', outputCol='pca_features')
pca_model = pca.fit(feature_df)
pca_df = pca_model.transform(feature_df)
```

```
train_dfer, test_dfer = feature_df.randomSplit([0.8, 0.2])
```

```
# Train a decision tree classifier
dter = DecisionTreeClassifier(featuresCol='features', labelCol='label', maxDepth=5 )
dt_modeler = dt.fit(train_dfer)

predictionser = dt_modeler.transform(test_dfer)

# Evaluate the model
evaluatorer = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction", metricName="accuracy")
accuracyer = evaluatorer.evaluate(predictionser)
print(f"Test set accuracy = {accuracyer}")
```

III. Evaluate

Result on test dataset after feature selection + dimensionally reduction

```
Test set accuracy = 0.5053239255933291
```

Result on test dataset with all features

```
Test set accuracy = 0.5239704593816498
```

=> We get the result above 50% among 6 classes, which is not so bad on Decision Tree with max 5 height

=> Note that we should tuning the model futhur - Like Adjust max_height of the tree, the method of calculating Impurity . However, my local computer is not efficiently do that so I will skipp

IV. Deploy the model and use on real-time

It is important to save the model to future use:

- I save the model
- Then load if use later

Real-time weather forecasting

Now it's time to do that in real world, we will collect data from real-time, use spark to process them quickly, then use the model to predict what is the weather currently

We will:

- Call API to get current weather data from many cities , we call every 5 minutes
- Preprocessing them to forming dataframe that similar to train/test set (using Pyspark ofcourse)
- Use model to predict weather
- We will also use the true label to test the model again

See the code in: **all_api_in_one**

V. Summary and Future Work

So far, we have address the weather forecasting using ML and Big Data Tools

There're still a lot of things we can improve this:

- Another ML / DL model
- Model Tuning
- More data collecting
- Better data Analysis techniques (LDA, StandardScaler,...)
- Finetuning model overtime when new data comes

*My local computer doesn't